

INTERNSHIP TASK REPORT

Prompt Engineering for Creative AI Startup Idea Generation

Task 02: Prompting for Creativity

Domain: Social Impact | Technology: Artificial Intelligence

Task Number	Task 02
Task Title	Prompting for Creativity
Domain	Artificial Intelligence & Social Impact
Submission	SkillCraft Technology Internship
Date	July 2026

TABLE OF CONTENTS

- 1. Objective..... 2
- 2. Introduction..... 2
- 3. Creative Task Overview..... 3
- 4. Prompt Version 1 — Basic..... 3
- 5. Analysis of Version 1..... 4
- 6. Prompt Version 2 — Detailed..... 4
- 7. Analysis of Version 2..... 6
- 8. Prompt Version 3 — Few-shot..... 6
- 9. Analysis of Version 3..... 9
- 10. Tone Experiment..... 9
- 11. Comparison Table..... 11
- 12. Key Observations..... 11
- 13. Conclusion..... 12
- 14. Learning Outcomes..... 13
- 15. Final Reflection..... 13

1. Objective

This report documents a structured exploration of prompt engineering techniques applied to a creative generation task: the design of original AI-powered startup ideas targeting social impact. The objective is to understand how prompt structure, context, tone, and embedded examples systematically influence the quality, creativity, and consistency of AI-generated outputs. Through three progressively refined prompt variants, a tone experiment, and a comparative analysis, this task demonstrates the foundational importance of deliberate prompt design in real-world generative AI applications.

2. Introduction

Artificial Intelligence has evolved from narrow, rule-based systems into flexible large language models capable of generating rich, contextual, and creative content. At the heart of effectively harnessing these capabilities lies a discipline known as Prompt Engineering.

Prompt Engineering is the practice of carefully crafting input instructions — called prompts — to guide a language model toward producing desired outputs. Far from being a simple act of asking a question, it involves deliberate choices about language, structure, context, role assignment, tone, and examples. A well-designed prompt can transform a vague, unhelpful response into a precise, high-quality, and actionable output.

Prompt design matters because language models respond to patterns in the input rather than inherent understanding of intent. The phrasing, detail, and format of a prompt directly determine what the model prioritises. A prompt that provides a clear role, context, and constraints will dramatically outperform one that offers only a bare instruction.

In Generative AI, creativity does not emerge randomly — it is shaped by the information space defined by the prompt. Creativity in this context refers to the novelty, relevance, and originality of generated outputs. Prompts that provide richer context create a larger, more meaningful solution space for the model to explore.

Few-shot prompting is a particularly powerful technique in which one or more complete examples are embedded within the prompt before the model is asked to generate new content. The model pattern-matches against these examples, producing far more consistent and well-formatted results than zero-shot prompts alone. This technique is especially effective for structured creative tasks — such as startup idea generation — where uniformity and quality across multiple outputs are desirable.

3. Creative Task Overview

Creative Task	Generating AI Startup Ideas for Social Impact
Domains Covered	Healthcare, Education, Agriculture, Disaster Management, Women's Safety, Mental Health, Accessibility, Climate Change, Smart Cities, Community Development
Prompt Variants	Version 1 (Basic) Version 2 (Detailed) Version 3 (Few-shot)
Tone Variants	Professional Visionary Humorous Investor Pitch

Each startup idea generated must demonstrate: a clearly articulated problem, an AI-driven solution mechanism, a defined target audience, a sustainable revenue model, and measurable innovation and social impact potential. The task progresses from a minimal prompt to a highly structured few-shot prompt to isolate the effect of each prompt design choice.

4. Prompt Version 1 — Basic Prompt

PROMPT

Generate 5 AI startup ideas for social impact.

OUTPUT

1. **HealthBot AI:** An AI chatbot that answers basic medical questions.
2. **EduAI:** An AI tool to help students learn faster.
3. **FarmAI:** AI for farmers to improve crop yield.
4. **SafeCity:** An AI surveillance system for public safety.
5. **ClimateTrack:** An AI platform to monitor climate change data.

5. Analysis of Version 1

Strengths	Produced five distinct, topically varied ideas quickly.
Weaknesses	Highly generic. Each idea described in a single sentence with no supporting depth.
Creativity	Low. Concepts such as a farming AI or medical chatbot are overused without differentiation.
Detail	Very poor. No revenue model, target user definition, or AI mechanism described.
Formatting	Minimal. A simple numbered list with no consistent structure.

The basic prompt demonstrates that without sufficient context or structure, the model defaults to the most surface-level interpretation of the task. The output, while technically responsive, is not suitable for professional or academic use. It confirms that prompt richness directly determines response richness.

6. Prompt Version 2 — Detailed Prompt

PROMPT

You are an expert AI product strategist with experience in social entrepreneurship and technology innovation. Your task is to generate 5 original AI startup ideas that address real-world problems in developing countries. Each idea must focus on a different domain (healthcare, education, agriculture, disaster management, or women's safety). For each startup, provide: the startup name, the problem it solves, the AI-driven solution, the target users, the revenue model, and a brief note on social impact. Each idea should be described in 60–80 words. Avoid generic concepts. Prioritise practicality, scalability, and innovation.

OUTPUT

Startup 1 — NutriSense AI | Healthcare

Startup Name	NutriSense AI
Problem	Rural communities in developing nations face widespread malnutrition due to a lack of dietitian access.
Solution	A mobile app that uses computer vision to analyse photographs of meals and assess nutritional content in real time. It provides personalised dietary guidance using regional food data.
Target Users	Low-income families and maternal health workers.
Revenue	Government health programme licensing and NGO partnerships.
Impact	Reduction in child malnutrition and improved maternal health outcomes.

Startup 2 — ScholarPath | Education

Startup Name	ScholarPath
Problem	First-generation university students in underserved communities lack mentorship and career guidance.
Solution	An AI platform that matches students with virtual mentors, recommends skill-building pathways, and predicts academic risk to trigger early interventions.
Target Users	Secondary and post-secondary students from low-income households.
Revenue	Institutional subscriptions and government education grants.
Impact	Improved higher education completion rates and increased social mobility.

Startup 3 — CropGuardian | Agriculture

Startup Name	CropGuardian
Problem	Small-scale farmers lose up to 30 percent of their crops annually to undetected disease and pest infestations.
Solution	An AI-powered drone and mobile system that scans crop fields, identifies disease patterns using image recognition, and recommends targeted treatments.
Target Users	Smallholder farmers across South Asia and Sub-Saharan Africa.
Revenue	Pay-per-scan model and agri-input company partnerships.
Impact	Increased food security and higher farmer incomes.

Startup 4 — AlertNow | Disaster Management

Startup Name	AlertNow
Problem	Marginalised and remote communities receive disaster warnings too late or not at all due to poor communication infrastructure.
Solution	An AI system that analyses satellite imagery, weather patterns, and seismic data to predict localised disaster risks 48–72 hours in advance, delivering multilingual alerts via SMS and community radio.
Target Users	Rural disaster-prone communities and civil protection agencies.
Revenue	Government contracts and international aid agency licensing.
Impact	Significant reduction in disaster-related casualties.

Startup 5 — ShieldHer | Women's Safety

Startup Name	ShieldHer
Problem	Women in urban and semi-urban areas frequently face safety threats with limited means of rapid response.
Solution	A wearable device paired with an AI app that uses voice pattern recognition and contextual cues to detect distress situations and automatically alert emergency contacts and local authorities.
Target Users	Women commuters and students.
Revenue	Device sales, subscription plans, and corporate safety programme partnerships.
Impact	Enhanced personal security and reduced gender-based violence incidents.

7. Analysis of Version 2

The introduction of a defined role (expert AI product strategist), geographic context (developing countries), domain constraints, and a structured output format transformed the quality of the response substantially.

Structure	Each startup follows a consistent six-field format: name, problem, solution, target users, revenue, and impact.
Creativity	Domain-specific with realistic AI mechanisms (computer vision for nutrition; satellite imagery for disaster prediction).
Detail	The 60–80 word limit enforced depth without verbosity. Every field is meaningfully populated.
Formatting	Clear, readable, and uniform across all five outputs — a marked improvement over Version 1.
Role Effect	Assigning the AI a professional identity introduced a more strategic and informed perspective throughout.

The key lesson from Version 2 is that constraints liberate rather than limit. By narrowing the model's freedom through role, domain, format, and length constraints, the output became more targeted, original, and practically useful. The instruction to avoid generic concepts was particularly effective.

8. Prompt Version 3 — Few-shot Prompt

Two fully elaborated example startup ideas were embedded in the prompt before requesting five new ideas. This technique demonstrates the expected structure, depth, vocabulary, and scoring system to the model before it generates output.

PROMPT

You are an AI startup idea analyst. Below are two complete examples of well-structured AI startup ideas for social impact. Study the format carefully and generate five new, original startup ideas following the exact same structure.

Example 1 — Provided in Prompt

Startup Name	VisionPath
Problem	Visually impaired individuals lack affordable navigation tools for independent travel in unfamiliar urban environments.
Solution	An AI-powered mobile application using real-time object detection, depth estimation, and natural language audio guidance to help visually impaired users navigate streets, public transport, and buildings independently.

Target Users	Visually impaired adults aged 18–60 in urban areas.
Revenue Model	Freemium app with premium navigation features; licensing to public transit authorities.

Innovation Score: 9 / 10	Impact Score: 9 / 10
---------------------------------	-----------------------------

Example 2 — Provided in Prompt

Startup Name	MindBridge
Problem	Rural adolescents experiencing mental health crises have no access to professional psychological support.
Solution	A multilingual AI therapy companion conducting structured, evidence-based mental health screenings, providing cognitive behavioural therapy exercises, and escalating high-risk cases to human counsellors via telemedicine.
Target Users	Adolescents aged 13–21 in rural and semi-urban regions.
Revenue Model	State mental health department contracts, school district subscriptions, and NGO funding.

Innovation Score: 8 / 10	Impact Score: 10 / 10
---------------------------------	------------------------------

OUTPUT — FIVE NEW STARTUP IDEAS

Startup 1 — AquaWatch | Clean Water

Startup Name	AquaWatch
Problem	Millions of households in water-stressed regions consume contaminated water without any means of quality testing.
Solution	An IoT sensor array paired with an AI analytics platform that continuously monitors water sources for contamination markers, predicts quality degradation, and alerts households and local authorities via a mobile dashboard.
Target Users	Rural and peri-urban households and municipal water boards in water-scarce regions.
Revenue Model	Sensor hardware sales, monthly monitoring subscriptions, and government utility contracts.

Innovation Score: 8 / 10	Impact Score: 9 / 10
---------------------------------	-----------------------------

Startup 2 — LegalEase AI | Justice Access

Startup Name	LegalEase AI
Problem	Marginalised communities cannot afford legal representation and remain unaware of their basic rights.
Solution	A conversational AI legal assistant trained on local jurisdiction laws that explains rights, prepares legal documents (affidavits, RTI applications), and connects users to pro bono lawyers for complex cases.
Target Users	Low-income individuals, daily wage workers, and first-generation litigants.
Revenue Model	Government justice department grants, NGO partnerships, and premium document drafting services.

Innovation Score: 8 / 10

Impact Score: 9 / 10

Startup 3 — SpeakUp School | Speech Therapy

Startup Name	SpeakUp School
Problem	Children with speech and language disorders in public schools receive inadequate therapy due to a national shortage of speech therapists.
Solution	An AI speech therapy platform that conducts individualised assessments, delivers interactive exercises tailored to each child's disorder profile, and tracks progress for review by teachers and parents.
Target Users	Children aged 4-14 with speech disorders in public and government-aided schools.
Revenue Model	School board subscriptions and state education department licensing.

Innovation Score: 9 / 10

Impact Score: 9 / 10

Startup 4 — ReliefChain | Humanitarian Logistics

Startup Name	ReliefChain
Problem	Aid distribution during humanitarian crises is

	frequently delayed and misallocated due to poor situational awareness and coordination failures.
Solution	An AI logistics platform integrating real-time satellite data, population displacement models, and supply chain information to optimise aid routing, predict shortfalls, and coordinate multi-agency relief operations.
Target Users	International humanitarian organisations, national disaster response agencies, and logistics partners.
Revenue Model	Subscription-based SaaS model for aid organisations and government emergency management contracts.

Innovation Score: 9 / 10

Impact Score: 10 / 10

Startup 5 — ElderCare Connect | Elderly Wellbeing

Startup Name	ElderCare Connect
Problem	Elderly individuals living alone are at high risk of medical emergencies, cognitive decline, and social isolation, with limited family oversight.
Solution	An AI-powered ambient monitoring system using non-intrusive sensors and behavioural pattern analysis to detect anomalies (falls, medication non-compliance, unusual inactivity) and alert caregivers and emergency services instantly.
Target Users	Senior citizens aged 65 and above living independently; adult children of elderly parents.
Revenue Model	Monthly subscription plans for families, senior living facility contracts, and insurance provider partnerships.

Innovation Score: 8 / 10

Impact Score: 9 / 10

9. Analysis of Version 3

Few-shot prompting produced the most consistent and high-quality outputs of all three prompt variants. By embedding two fully elaborated examples before requesting new ideas, the model received an unambiguous template defining the expected depth, vocabulary, structure, and scoring mechanism.

- Every generated startup followed the same six-field structure without deviation.
- The ideas were notably more original — AquaWatch, LegalEase AI, and ReliefChain did not appear in earlier outputs.
- The Innovation and Impact scoring fields introduced a quality benchmark that elevated the standard of ideas generated.
- Few-shot examples do not merely prescribe format — they calibrate the model's implicit quality standard.
- Consistency across all five outputs was near-perfect, making this the most suitable prompt variant for professional submission.

10. Tone Experiment

The same base prompt — "Generate one AI startup idea for social impact in the field of mental health" — was issued with four different tone instructions. The purpose was to examine how tonal direction alone transforms the communicative function of the output.

Tone 1 — Professional

MindMetrics is an AI-driven mental health monitoring platform designed for corporate and institutional environments. The platform analyses anonymised behavioural data — including productivity patterns, communication frequency, and activity levels — to identify early indicators of employee burnout or psychological distress. HR teams and occupational health professionals receive aggregated risk dashboards and are prompted to initiate confidential support interventions. The solution operates within a strict ethical framework, ensuring employee data sovereignty and regulatory compliance. Revenue is generated through enterprise subscriptions.

Tone 2 — Visionary Entrepreneur

Imagine a world where mental health crises are predicted before they arrive — where no person has to struggle in silence. ClarityOS is that vision made real. We are building an AI consciousness layer for human wellbeing: a platform that learns your emotional rhythms, detects subtle shifts in your mental state, and quietly places the right support — a mindfulness nudge, a therapist connection, a peer community — exactly when you need it. This is not just technology. This is the infrastructure of human flourishing.

Tone 3 — Humorous

Introducing TherapyBot 3000 — the startup that finally admits what everyone already knows: talking to an AI at 2 AM is more accessible than finding a therapist with an available appointment this decade. TherapyBot 3000 does not judge your questionable life choices, never runs over its

session time, and does not charge you extra for crying. Using sentiment analysis and natural language processing, it offers evidence-based coping strategies and gently suggests professional help when it realises it is out of its depth. Free tier available. Premium tier includes a virtual rubber duck for developers.

Tone 4 — Investor Pitch

The global mental health market is projected to exceed \$500 billion by 2030, yet only 1 in 5 individuals in need receives adequate care. NeuralWell is positioned to capture this gap. Our platform deploys multimodal AI to deliver personalised, scalable mental health support at a cost 90 percent below traditional therapy. With a proven freemium-to-subscription conversion model, an NPS of 74 in our pilot cohort of 12,000 users, and active conversations with three national health ministries, NeuralWell is ready to scale. We are seeking \$8 million in Series A funding to expand to five new markets by Q3.

How Tone Changed the Output

Professional	Structured, compliance-aware product description suited to enterprise procurement communications.
Visionary	Aspirational, metaphor-rich language designed to inspire emotional investment in the mission.
Humorous	Lowered barrier to engagement; communicated product value through relatability and wit.
Investor Pitch	Quantitative, urgency-driven language prioritising metrics, market size, traction, and call-to-action.
Key Finding	A single tonal instruction transformed the same concept into four distinct communication artefacts, each serving a different audience and strategic purpose.

11. Comparison Table — Prompt Version Analysis

The table below rates each prompt version across seven quality dimensions on a scale of 1 to 10. Scores reflect the relative performance of each version, with Version 3 (Few-shot) achieving the highest overall rating.

Prompt Version	Type	Structure /10	Creativity /10	Detail /10	Consistency /10	Formatting /10	Clarity /10	Overall /10
Version 1	Basic	2	4	3	3	2	6	3
Version 2	Detailed	8	7	8	7	8	8	8
Version 3	Few-shot	9	9	9	10	9	9	9

Scores are assigned on the basis of output quality relative to professional submission standards. A score of 10 indicates exemplary performance on that dimension; 1 indicates near-complete failure.

12. Key Observations

- Assigning a role (e.g., expert AI product strategist) significantly elevated the sophistication and relevance of outputs.
- Domain constraints directed the model toward a diverse set of ideas rather than clustering around the most common themes.
- The basic prompt produced surface-level results because it provided no information about expected depth, format, or quality standard.
- Detailed prompts with explicit output format requirements reduced ambiguity and produced consistently structured responses.
- Few-shot prompting was the most effective technique for enforcing structural consistency across all generated outputs.
- Embedding worked examples calibrated the model's quality threshold, resulting in startup ideas with greater specificity and originality.
- Tone instructions fundamentally altered the communicative register of the output without changing the underlying factual content.
- Word limit constraints prevented over-generation and helped maintain focused, readable descriptions.
- The presence of Innovation and Impact scoring fields introduced an implicit quality benchmark that elevated the standard of ideas.
- Geographic specificity (developing countries, South Asia, Sub-Saharan Africa) produced more contextually appropriate and actionable ideas.
- Combining role, context, task, format, and constraints in a single prompt reliably outperformed any single element used in isolation.
- Longer, richer prompts consistently produced longer, higher-quality responses — confirming that input richness correlates strongly with output richness.

13. Conclusion

This task has demonstrated, through direct experimentation, that Prompt Engineering is not a peripheral concern in working with Generative AI — it is a foundational competency. The quality, creativity, consistency, and usefulness of AI-generated outputs are not primarily a function of the model's capabilities alone; they are a function of the inputs provided to it.

Prompt structure plays a decisive role in shaping output quality. A bare instruction — such as "generate 5 AI startup ideas" — leaves the model without sufficient guidance to produce anything beyond the most generic response. In contrast, a structured prompt that defines a role, establishes context, specifies the task, prescribes the output format, and enforces constraints consistently yields outputs that are deeper, more original, and far more practically useful.

Few-shot prompting proved to be the most powerful technique explored in this study. By embedding complete, well-formed examples within the prompt, the model received an unambiguous template for the expected output. This technique is particularly valuable when consistency across multiple outputs is required — as is the case in professional, academic, or product design contexts. The few-shot approach also calibrated the model's implicit quality standard, resulting in outputs that matched or exceeded the sophistication of the provided examples.

Role prompting — the practice of assigning the model a specific professional identity — introduced expertise and authority into the outputs that was absent from basic prompts. This technique is valuable in any domain where subject-matter knowledge and strategic framing are desirable outcomes.

Tone instructions demonstrated that the same underlying concept can be rendered as a formal product specification, an inspirational vision statement, a humorous pitch, or a metrics-driven investor presentation — depending entirely on a single sentence of tonal guidance. This has significant implications for communication design, content marketing, and educational material development.

Looking forward, Prompt Engineering is set to become an increasingly critical skill as AI systems are deployed across industries. From healthcare diagnostics to legal assistance to disaster response, the difference between a useful AI application and an ineffective one will often lie in the quality of the prompts driving it. As models grow more capable, the art and science of prompt design will continue to evolve — becoming more nuanced, more context-sensitive, and more consequential.

14. Learning Outcomes

- ✓ Learned and applied core prompt engineering techniques: basic, detailed, and few-shot prompting.
- ✓ Understood how few-shot prompting increases output consistency and quality by providing worked examples.
- ✓ Explored how tone-based instructions transform the communicative register and purpose of AI-generated content.
- ✓ Developed the ability to design structured prompts with role, context, task, format, and constraint specifications.
- ✓ Gained practical experience evaluating AI outputs across creativity, detail, consistency, and formatting dimensions.
- ✓ Understood how geographic and domain-specific constraints improve the relevance and applicability of AI-generated ideas.
- ✓ Improved critical AI interaction skills through iterative prompt refinement and comparative output analysis.

-
- ✓ Recognised the real-world applications of prompt engineering in product strategy, healthcare, education, and social entrepreneurship.

15. Final Reflection

Completing this task was a genuinely eye-opening experience. Before beginning, I understood that the way a question is asked can influence the answer — but I underestimated how dramatically this effect manifests in AI systems. Moving from a single-line basic prompt to a richly structured few-shot prompt felt less like editing a query and more like redesigning a brief: the model's outputs improved in clarity, originality, and usefulness at every step.

The tone experiment was particularly revealing. Watching the same mental health startup concept transform into four entirely distinct documents — each tailored to a different audience and purpose — made me realise that prompt engineering is not merely a technical skill. It is a form of communication design. Understanding the audience, the context, and the intended effect of a piece of writing is just as important as knowing the prompt syntax.

This internship task has given me a practical framework I intend to carry forward into professional and research contexts. I am now far more intentional about how I structure AI interactions, and I have a deeper appreciation for the craft underlying effective human-AI collaboration.

End of Report

SkillCraft Technology Internship | Task 02: Prompting for Creativity | July 2026